

Recommendation Systems Approaches on the Netflix Prize Data Set

Cody Blakeney
Email: cjb92@txstate.edu

Samuel Teich
Email: st1289@txstate.edu

Abstract—Recommendation systems play an increasingly vital role in the ability to make informed decisions in modern societies, given the ever growing deluge of data which might inform those decisions. Without the ability to quickly choose from among a few well selected options ordinary activities from applying to jobs to searching the internet, buying household goods, or reading the latest news might become difficult, deceptive, or even dangerous. Given the many relatively recent major advances in machine learning techniques, especially in relation to deep learning, it is important to understand how recommendation systems have changed and improved. In this paper we investigate approaches to recommendation systems using deep learning.

1. Introduction

Data characterizes the modern world. Corresponding to exponential growth in computing power over time, and the related proliferation of sensors and interactions with computing systems, the volume of recorded data has exploded. It is estimated that in 2010 there was more than 1 Zettabyte of recorded unstructured data, and that as of 2018 there exists more than 2 Zettabytes structured data and 20 Zettabytes structured data [2][8]. While much of this deluge is comprised of highly specific and narrowly specialized data, not useful for making decisions at a personal level, there still is a wealth of data - far more than is human interpretable - which could be used to inform such decisions. Recommendation systems are predictive tools designed specifically for this use case, helping pare from a glut of options a select few choices based on available data.

The Netflix Prize competition [1], started in 2006 and finalized in 2009, is, along with the publishing of the ImageNet dataset in 2009 [3], credited as one of the events that encouraged the perception of a revolution in the field of machine learning both among professional audiences and the public at large [7]. In the near decade since team BellKor's Pragmatic Chaos won the Netflix Prize competition [1], rapid advances, especially in Deep Learning, have continued to push boundaries in machine learning and recommendation systems. The Netflix dataset remains an important benchmark dataset for recommendation systems.

In this paper we will discuss recent improvements in the field of recommendation systems and attempt to benchmark some of those improvements on the Netflix dataset. Section 2 will cover some important background material on the

field of recommendation systems. Section 3 will discuss some recent and related results. Section 4 will talk about our experimental setup, while Section 5 will include the specific experiments carried out. Section 6 discusses our results, conclusions, and will cover our expectations for future work.

2. Background

2.1. Recommendation Systems

Recommendation systems predict the preferences of users in regard to new items or objects, in the form of a ranking of items or objects, rating items or objects, or classifying items or objects [11]. There exist a wide variety of established strategies to learn prediction models for recommendation systems, generally fitting into three categories: collaborative filtering, content based, or hybrid systems. Collaborative filtering based recommendation system models learn based on user - item/object interaction, while content based recommendation system models learn based on intrinsic user and item/object characteristics. Hybrid recommendation system models learn based on a combination of the two approaches.

2.2. Deep Learning

Deep learning is a subset of machine learning wherein models learned contain multiple, often hierarchical, abstractions or representations. While there are many approaches to deep learning, in this paper we employ deep feedforward neural networks - feed forward neural networks with one or more hidden layers. A general interpretation of the hidden layers in a deep feedforward neural network is that each layer learns a slightly more complex representation of the data than the layer before it.

2.3. Embedding Layers

Embedding layers encode information about a set of objects and the relationships between those objects in a latent space - generally with fewer dimensions than the original space representing the objects [6]. For example, given a corpus of text with a large vocabulary V , it may be desirable to represent $v \in V$ such that $\|rep(v)\| = k$, where $k \ll |V|$ and that if $v, x \in V$ occur together more frequently than $v, y \in V$, v will be closer to x than y .




Figure 1. Simple Embedding Layer Architecture

2.3.1. Neural Network Embedding Layers. An embedding layer as part of a neural network is framed as the weights of the hidden layer in a single hidden layer dense feedforward neural network. Generally there are two ways of training such a model; by predicting the target object given that object's context as input (Continuous Bag of Words - CBOW) or by predicting the context of a target object given that object as input (Skip-gram).

		Item										
		W	X	Y	Z			W	X	Y	Z	
User	A		4.5	2.0		=	X	A	1.2	0.8		
	B	4.0		3.5				B	1.4	0.9		
	C		5.0		2.0			C	1.5	1.0		
	D		3.5	4.0	1.0			D	1.2	0.8		
Rating Matrix								User Matrix				
								Item Matrix				

Figure 2. MF diagram

2.3.2. SVD Embedding Layers. Singular Value Decomposition is a form of Matrix Factorization popular in tradition recommendation systems. The idea is to start with a matrix A where each entry represents an interaction. In the case of a User-movie embedding (here after referred to as user embedding), the rows of A represent a user and the columns represent a movie. As you can see from figure 2 Every entry is either the users rating of that movie, or zero. You then decompose your matrix to the product of three matrices $A = U\Sigma V^t$ A is an $M \times N$ matrix U is a $M \times M$ matrix Σ is an $M \times N$ rectangular diagonal matrix and V is a $N \times N$ unitary matrix. To get a reduced approximation of A we simply decide how many k features we want and keep the k most important features of Σ . Multiplying $U\Sigma$ will result in A of rank k . In traditional recommendation systems the multiplying the reduced rank Σ with all three matrices would result in a new approximation of A with values filled in from the reduced rank approximation. The predicted ratings are then looked up in the matrix for an item. For our use case we were only intrested in the reduced space representatoin of the matrix not transforming it back so we kept only $U\Sigma$

3. Related

The winning team of the Netflix Prize competition, team BellKor's Pragmatic Chaos, used a large ensemble of smaller models, including decision trees, support vector machines, restricted boltzmann machines, and regression models in their solution, but did not use any deep learning models - deep learning models were not yet popular or practical [?][10][5]. Since then, deep learning recommendation systems have been the subject of extensive research. A survey paper, [11], detailing several such deep learning recommendation systems, including models employing deep feedforward neural networks (MLP), recurrent neural networks (RNN), convolutional neural networks (CNN), deep auto-encoders (AE), generative adversarial networks (GAN), and composite models. Given project time restrictions, our we narrowed our focus to MLP and RNN based deep learning recommendation systems.[4] Details an approach using MLP and both neural network embedding layers and SVD embeddings. In [?] multiple data sources are used as disparate points in a heierarchical model to rank youtube videos, and the challenges of implementation at scale are discussed. Non-linear matrix factorization for sparse inputs with a collaborative filtering recommendation system model is discussed in [9], and state of the art results are given on the movieLens and jester datasets.

3.1. Takeaway From Related Work

There were four main themes throughout the body of related work:

- Meaningful object representation
- Explicit hierarchical abstraction
- Control of data flow
- Use of domain or expert knowledge

Black box solutions - such solving the problem purely through application of large, overparameterized, neural networks - do not appear in the literature on deep learning recommendation systems in the same way that they seem to in the fields of image segmentation or recognition. Deep learning recommendation systems seem to require to be tailored for the exact context and medium of the interaction. As such methods from implicit and explicit interaction have somewhat limited interoperability. It seems the trend in recommendations is moving towards implicit interactions which, given our dataset limits some of what we can do. Given time constraints, we chose to focus our efforts on meaningful object representation.

4. Approach

Our approach to investigating the architecture for deep learning recommendation systems, was to first decide on a baseline with which we could compare the effectiveness of our techniques. After finding a benchmark we reviewed the literature we had assembled to decide which design was most applicable to our dataset, and our ability to implement.

We knew from the Netflix competition that the winners had an RMSE of 0.8567. We also wanted a way to compare our engineered solution, to one that simply took all the data Netflix provided and used it as inputs on a naive feed-forward neural net.

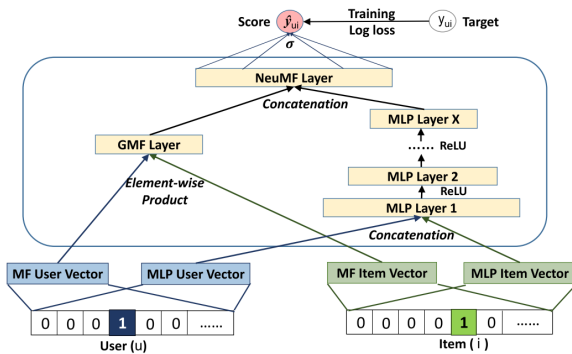


Figure 3. NCF diagram

The theme of all the papers we read covering the various methodologies for applying deep learning to recommendation systems, was creating more meaningful representations of the users, the items they interacted with, and context around the item. Many of the models we read about involved many complex hierarchical layers. While many of these exotic arrangements made sense, we were not sure we had the technical expertise to implement them in the allotted time. We settled on a technique called Neural Collaborative Filtering [4]. Neural Collaborative Filtering or NCF is a process that involves creating a hybrid of traditional recommendation system and deep learning. As you can see from figure 3 it combines user vectors and item vectors from matrix factorization as well as user and item embedding layers as inputs to the neural network.

In order to try to replicate this architecture we would need to need to find appropriate Matrix Factorization representations for our users and our movies, as well as create embedding layers for for them.

5. Setup

5.1. The Netflix Prize Dataset

The Netflix dataset is formatted as a set of 100480507 rating events, each event including the user, movie, and date corresponding to the rating and the 1 - 5 score of the rating itself. In all, the dataset includes ratings from 480189 anonymized users in reference to 17770 individual movies and tv-show episodes.

5.2. The IMDB Dataset

The IMDB dataset is refreshed daily and is comprised of seven separate sub-datasets, of which only three were used:

- title.basics.tsv - Basic information on titles in the IMDB catalog
- title.principals.tsv - Principal cast and crew for each title in the IMDB catalog
- name.basics.tsv - Basic information for cast and crew

In total, 2290948 different titles for movies and tv-show episodes are included in the dataset and 8469380 different names corresponding to cast and crew.

5.3. Experimental Setup

Experiments were conducted using Tensorflow-gpu 1.8.0, Keras 2.1.4, and CUDA 9.0. The machine on which the experiments ran had 32g ram, an i7-6700K CPU, and a GeForce GTX 1080 GPU.

6. Experiments

6.1. Naive Feed-forward Deep Neural Networks

What we call our Naive Feed-forward network is the benchmark we made to test the efficacy of the models we were training. We tested several versions before deciding on a final candidate Naive Model. We had a categorical classifier, a regression model, and another categorical model with a better layer design.

6.1.1. Naive Categorical Model. Our first model had inputs of one hots for movies and users, mean, median and variance of the current movie calculated from the training set, and time represented as a month input, a day input, and a year since beginning input. This model had two hidden layers with 128 nodes each using relu activation and a softmax output to guess categorically what the rating was for that user. With this configuration the model converged to about 38% accuracy on our validation set. Because of the distribution of ratings in our data set this is about as good as always guessing the rating 4 for every movie. Not great, but better than randomly chance.

6.1.2. Naive Regression Model. For training the problem as a regression model we learned we had to be very careful with the layer selection. Too few layers and inadequate funneling down of the layers towards the output layer resulted in a model that was hopelessly lost at trying to find a gradient. We finally had success with a model using a hidden layers of size 512, 64, 8, and a linear output layer. After training we were surprised because it had a very good RMSE score. Closer inspection showed that the model was predicting a floating point value that was always between 3.4 and 3.6, very close to the mean value of the data set. When we rounded the value and calculated the accuracy and RMSE it was only 21% accurate with and RMSE of 1.203.

6.2. Neural Network Embeddings on the Netflix Dataset

Since it was unclear how, or whether, to attempt to model ratings in the latent space, several experiments were carried out on both the movie embeddings and user embeddings over the Netflix dataset:

- User Embeddings trained as CBOW across all rating events
- User Embeddings trained as CBOW across positive (≥ 3 rating) rating events
- User Embeddings computed as a weighted mean of all movie embeddings for movies that a user has watched
- Movie Embeddings trained as CBOW across all rating events
- Movie Embeddings trained as CBOW across positive (≥ 3 rating) rating events
- Movie Embeddings computed as a weighted mean of all user embeddings for movies that a user has watched

Latent Spaces for all experiments were 64 dimensions, and all embedding layers were trained with noise contrastive error loss. Weights were calculated as -1 for ratings of 1 and 2, and 1 for ratings of 3, 4, and 5 to attempt to separate movies in the latent space by rating sentiment. All embedding layers were trained for at least 100 epochs. Unfortunately an objective measure of the quality of a latent space was lacking in this case; NCE loss does not necessarily reflect the quality of the representation of semantic (object / object in this case) relationships across the latent space. To measure quality random samples were visualized across the latent space using TSNE and clustering was manually verified.

6.3. Neural Network Embeddings on the IMDB Dataset

A principle contributor neural network embedding layer was built using CBOW training - treating the context for a principle as the other principles that had worked in conjunction with the target on a movie. The embedding dimension was 64, and this model was trained using noise contrastive error as loss for 500 epochs to a loss of 5.23. Movies embeddings were then computed as the mean of their constituent principle embeddings. Since the Netflix and IMDB movie catalogs had significant overlap, but neither were a subset of the other, multitask training was planned to form a joint embedding space merging the features of both embeddings. In this case a new latent space would be by auto-encoder from the existing IMDB and Netflix embedding layers: for samples in one set but not the other the loss would be computed only on the set with membership, but for samples in both movie databases the loss would be computed jointly. Unfortunately due to time constraints, this could not be finished.




Figure 4. TSNE visualization of unweighted embeddings over all movies for 100 movies in the Netflix movie catalog




Figure 5. TSNE visualization of weighted embeddings over all movies for 100 movies in the Netflix movie catalog

6.4. SVD Embeddings on the Netflix Dataset

In [4] they use a process the call GMF or Generalized Matrix Factorization where they use gradient descent to train their embedding to learn the most important latent features

When considering the methods for representing User and Movie matrix factorization embeddings, we decided on using Single Value Decomposition, primarily because it is implemented in the scipy package, and could operate on




Figure 6. TSNE visualization of embeddings for 500 movies in the IMDB movie catalog

sparse matrices. This was important because the original matrix we would be decomposing would be of size 440,000 x 17776 and would not possibly fit into memory.

6.5. Deep Feed-forward Neural Networks with Embedding Layers

A series of deep feed forward networks were trained using pairs of the netflix user / movie embedded layers as input: in each case these networks were fully connected, used leaky-relu activations with an alpha of .3 for the hidden layers and softmax for the output, had hidden layers of shape 512-128-32, and were optimized by SGD. For validation, 10% of the training data was held out. Loss was categorical crossentropy, and each network was trained as a classification problem predicting probability of target rating 1-5. Unfortunately, all of these experiments failed, results are not reported. In no case did validation accuracy ever rise, and in some cases validation accuracy fell.

6.5.1. SVD Deep Embeddings. We needed some way to evaluate how useful the information presented by the SVD Embeddings would be to a deep learning model. We decided to train models on combinations of embeddings and one hots only, without other contextual information like, date of review, actor, genre, etc to see how useful these representations would be. This would tell us, given only a representation of likeness to other users we had computed, how well we could predict their tastes. If we could beat the threshold of randomly guessing or always rating 4, the marks set by our benchmark naive model, than the matrix factorization embeddings would be considered a success.

To test this we first trained to convergence a model that took as inputs both the User SVD embedding and a one hot for its movie id. It was during training this model we experimented with adding more layers and funneling their output as mentioned in [4] with hidden layers of our input size (17826), 1000, 100, 10, and a softmax output layer we were able to achieve an accuracy of around 42%

We also tested using both the User SVD and the Movie SVD as the inputs. This would shrink the input size down from 17826 to 125. Training a model that again had 4 hidden layers we were able to reach an accuracy of 45%. Clearly the representations did have value.

Unfortunately we were not able to also train a model using user one hots and movie embeddings, as the input size made computation prohibitively difficult.

6.6. Hierarchical Architectures

As mentioned in our the approach section, our original intention was to construct our own version of a model similar to that of NCF as in figure 3. Because of setbacks finding powerful embedding representation of user item interactions we were unable to construct that Hierarchical.

Even though we do not have experimental results from these models we would like to talk a bit about what exactly we were planning to build to replicated, and extend [4] and what our reasoning was. As you can see in figure 3 in NCF embedding layers are run through several MLP layers, while the GMF their matrix factorization is simply found and concatenated at a joint layer just before the output layer. Given our positive results training a network on just the output of matrix factorization we reasoned that more useful arrangements could be made. We proposed the following two architectures.

- DNN SVD concatenated with DNN embedding layers
- combined SVD and embedding layer inputs

For the first model we would simple take our two trained models, the SVD Deep Embeddings model and the Deep model based on Netflix embeddings, remove there output layers and use them as inputs into a combining layer before a final output layer. This would allow learned information from different contexts to be combined in non-linear ways and possibly produce more accurate results.

The second model we propose is simply to train one wide and deep model where the inputs are both the SVD embedding layers and the Netflix embedding layers. The difference being that it would learn for itself the relationship between the embeddings and decide how to combine them.

7. Conclusions

While we have had mixed results with some of our experiments its clear that blindly entering data to a feed-forward DNN will not give you good results for recommendation engines.

A deep matrix trained on just the reduced space of user item interactions can perform reasonably well even without the dimensionality of those spaces being optimized, and being provided no other contextual information about the user or movie.

7.1. Experimental Discussion

The results of our experiments with the neural network derived embedding layers feeding into a deep feedforward network were disheartening. While it is possible that simply training longer, or with a slightly different architecture, optimizer, different activations, or regularization would improve performance, these negative results likely mean that the features of the latent space learned from the neural network derived embedding layers were not meaningful. Of course, this begs the question of what went wrong - but there's a larger lesson to be learned in how the quality of an embedding is measured. In our case measuring the quality of semantic relationships on these embeddings had to be measured manually and TSNE was our best visualization tool to see if similar movies or users had actually ended up clustered together in the latent space. Unfortunately, it seems this is an unreliable way to assessing quality, and we may have told ourselves a just-so story. In the future, better methods must be developed or researched.

7.2. Future Work

There is still much to be done as limited time and a steep learning curve prevented us from accomplishing some of what we set out to do. It is clear that our Matrix factorization DNNs were able to learn something more from the representations we provided than the raw data. We would like to find some way to fine tune the number of features we select in the resulting embedding space to maximize its usefulness in the future. We may explore some method like GFM to accomplish this. We also need to work to find more useful neural network embeddings for our users and items.

While our work in this paper focused mainly on creating better representations for user movie interactions, there are lots of rich forms of data we have left unexplored that could help our model. providing information about genre, sentiment representation of descriptions, volume of ratings during the user movie interaction period (as a measure of popularity) could help provide models with significantly more information about the movies. Additional creating some representation for the viewing patterns of users could also enrich the model. These are all areas we have left untouched that each require extensive investigation. As we have learned in our experiments to date, models are very sensitive to how data is represented to them.

In addition to expanding our methods of creating better representations and providing richer context, we could also explore different network methods like using auto encoders. They could be used both as a replacement for our embedding layers, using the bottle neck layer, or having the auto encoder fill blacks in the reconstruction layer.

References

- [1] [Online]. Available: <https://www.netflixprize.com/index.html>
- [2] J. R. David Reinsel, John Gantz, "Data age 2025: The evolution of data to life-critical," 04 2017. [Online]. Available: <https://www.seagate.com/files/www-content/our-story/trends/files/Seagate-WP-DataAge2025-March-2017.pdf>
- [3] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and F.-F. Li, "Imagenet: a large-scale hierarchical image database," pp. 248–255, 06 2009.
- [4] X. He, L. Liao, H. Zhang, L. Nie, X. Hu, and T.-S. Chua, "Neural Collaborative Filtering," 2017. [Online]. Available: <http://arxiv.org/abs/1708.05031>
- [5] Y. Koren, "1 the bellkor solution to the netflix grand prize," 2009.
- [6] T. Mikolov, K. Chen, G. Corrado, and J. Dean, "Efficient estimation of word representations in vector space," *CoRR*, vol. abs/1301.3781, 2013. [Online]. Available: <http://arxiv.org/abs/1301.3781>
- [7] H. Quentin, "Reasons to believe the a.i. boom is real," 03 2018. [Online]. Available: <https://www.nytimes.com/2016/07/19/technology/reasons-to-believe-the-ai-boom-is-real.html>
- [8] L. Rizzatti, "Digital data storage is undergoing mind-boggling growth," 09 2016. [Online]. Available: https://www.eetimes.com/author.asp?section_id=36&doc_id=1330462
- [9] F. Strub and J. Mary, "Collaborative Filtering with Stacked Denoising AutoEncoders and Sparse Inputs," in *NIPS Workshop on Machine Learning for eCommerce*, Montreal, Canada, Dec. 2015. [Online]. Available: <https://hal.inria.fr/hal-01256422>
- [10] A. Tscher and M. Jahrer, "The bigchaos solution to the netflix grand prize," 01 2009.
- [11] S. Zhang, L. Yao, and A. Sun, "Deep learning based recommender system: A survey and new perspectives," *CoRR*, vol. abs/1707.07435, 2017. [Online]. Available: <http://arxiv.org/abs/1707.07435>